

How you can drive safely in 2020

Use LR/KNN/Decision Tree/Random Forest classification algorithms from sklearn to predict the accident severity

Due to the limit of computer capacity, I am focusing on the State of California. I will only select a few features I believe are more relevant to severity. Categorical data will be treated with Pandas get_dummies method. Rows with missing values will be dropped.

Data source

<https://www.kaggle.com/sobhanmoosavi/us-accidents> (<https://www.kaggle.com/sobhanmoosavi/us-accidents>)

Acknowledgements

Moosavi, Sobhan, Mohammad Hossein Samavatian, Srinivasan Parthasarathy, and Rajiv Ramnath. "A Countrywide Traffic Accident Dataset.", 2019.

Moosavi, Sobhan, Mohammad Hossein Samavatian, Srinivasan Parthasarathy, Radu Teodorescu, and Rajiv Ramnath. "Accident Risk Prediction based on Heterogeneous Sparse Data: New Dataset and Insights." In proceedings of the 27th ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems, ACM, 2019.

Author

Ronghui Zhou, zhou.uf@gmail.com <https://github.com/RonghuiZhou> (<https://github.com/RonghuiZhou>)

Step 1. Import libraries

```
In [1]: # Import numpy, pandas, matplotlib.pyplot, sklearn modules and seaborn
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
pd.set_option('display.max_rows', 200)
pd.set_option('display.max_columns', 200)
plt.style.use('ggplot')

# Import KNeighborsClassifier from sklearn.neighbors
from sklearn.neighbors import KNeighborsClassifier

# Import DecisionTreeClassifier from sklearn.tree
from sklearn.tree import DecisionTreeClassifier

# Import RandomForestClassifier
from sklearn.ensemble import RandomForestClassifier

# Import LogisticRegression
from sklearn.linear_model import LogisticRegression

from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_curve, auc
```

Step 2. Import the dataset

```
In [2]: # Import the data  
df = pd.read_csv('./US_Accidents_May19.csv')  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2243939 entries, 0 to 2243938
Data columns (total 49 columns):
ID                object
Source            object
TMC               float64
Severity          int64
Start_Time        object
End_Time          object
Start_Lat         float64
Start_Lng         float64
End_Lat           float64
End_Lng           float64
Distance(mi)      float64
Description        object
Number            float64
Street            object
Side              object
City              object
County            object
State             object
Zipcode           object
Country           object
Timezone          object
Airport_Code      object
Weather_Timestamp object
Temperature(F)    float64
Wind_Chill(F)     float64
Humidity(%)       float64
Pressure(in)      float64
Visibility(mi)    float64
Wind_Direction    object
Wind_Speed(mph)   float64
Precipitation(in) float64
Weather_Condition object
Amenity           bool
Bump              bool
Crossing          bool
Give_Way          bool
Junction          bool
No_Exit           bool
Railway           bool
Roundabout        bool
Station           bool
Stop              bool
Traffic_Calming   bool
Traffic_Signal    bool
Turning_Loop      bool
Sunrise_Sunset    object
Civil_Twilight    object
Nautical_Twilight object
Astronomical_Twilight object
dtypes: bool(13), float64(14), int64(1), object(21)
memory usage: 644.1+ MB
```

Step 3. Extract year, month, day, hour, weekday, and time to clear accidents

```
In [3]: # Convert Start_Time and End_Time to datetimes
df['Start_Time'] = pd.to_datetime(df['Start_Time'], errors='coerce')
df['End_Time'] = pd.to_datetime(df['End_Time'], errors='coerce')

# Extract year, month, day, hour and weekday
df['Year']=df['Start_Time'].dt.year
df['Month']=df['Start_Time'].dt.strftime('%b')
df['Day']=df['Start_Time'].dt.day
df['Hour']=df['Start_Time'].dt.hour
df['Weekday']=df['Start_Time'].dt.strftime('%a')

# Extract the amount of time in the unit of minutes for each accident, round t
o the nearest integer
td='Time_Duration(min)'
df[td]=round((df['End_Time']-df['Start_Time'])/np.timedelta64(1,'m'))
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2243939 entries, 0 to 2243938
Data columns (total 55 columns):
ID                object
Source            object
TMC               float64
Severity          int64
Start_Time        datetime64[ns]
End_Time          datetime64[ns]
Start_Lat         float64
Start_Lng         float64
End_Lat           float64
End_Lng           float64
Distance(mi)      float64
Description        object
Number            float64
Street            object
Side              object
City              object
County            object
State             object
Zipcode           object
Country           object
Timezone          object
Airport_Code      object
Weather_Stamp     object
Temperature(F)    float64
Wind_Chill(F)     float64
Humidity(%)       float64
Pressure(in)      float64
Visibility(mi)    float64
Wind_Direction    object
Wind_Speed(mph)   float64
Precipitation(in) float64
Weather_Condition object
Amenity           bool
Bump              bool
Crossing          bool
Give_Way          bool
Junction          bool
No_Exit           bool
Railway           bool
Roundabout        bool
Station           bool
Stop              bool
Traffic_Calming   bool
Traffic_Signal    bool
Turning_Loop      bool
Sunrise_Sunset    object
Civil_Twilight    object
Nautical_Twilight object
Astronomical_Twilight object
Year              int64
Month             object
Day               int64
Hour              int64
Weekday           object

```

```
Time_Duration(min)      float64
dtypes: bool(13), datetime64[ns](2), float64(15), int64(4), object(21)
memory usage: 746.9+ MB
```

Step 4. Deal with outliers

A. Drop rows with negative time_duration

```
In [4]: # Check if there is any negative time_duration values
df[td][df[td]<=0]
```

```
Out[4]: 69720      -0.0
        69721      -0.0
        69722      -0.0
        69723      -1.0
        69724      -1.0
        309389    -30.0
        309390    -30.0
        746173    -30.0
        746174    -31.0
        1482940   -30.0
        1483025   -30.0
        1483026   -31.0
        2038017   -31.0
Name: Time_Duration(min), dtype: float64
```



```
In [5]: # Drop the rows with td<0

neg_outliers=df[td]<=0

# Set outliers to NAN
df[neg_outliers] = np.nan

# Drop rows with negative td
df.dropna(subset=[td],axis=0,inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2243926 entries, 0 to 2243938
Data columns (total 55 columns):
ID                object
Source            object
TMC               float64
Severity          float64
Start_Time        datetime64[ns]
End_Time          datetime64[ns]
Start_Lat         float64
Start_Lng         float64
End_Lat           float64
End_Lng           float64
Distance(mi)      float64
Description        object
Number            float64
Street            object
Side              object
City              object
County            object
State             object
Zipcode           object
Country           object
Timezone          object
Airport_Code      object
Weather_Timestamp object
Temperature(F)    float64
Wind_Chill(F)     float64
Humidity(%)       float64
Pressure(in)      float64
Visibility(mi)    float64
Wind_Direction    object
Wind_Speed(mph)   float64
Precipitation(in) float64
Weather_Condition object
Amenity           float64
Bump              float64
Crossing          float64
Give_Way          float64
Junction          float64
No_Exit           float64
Railway           float64
Roundabout        float64
Station           float64
Stop              float64
Traffic_Calming   float64
Traffic_Signal    float64
Turning_Loop      float64
Sunrise_Sunset    object
Civil_Twilight     object
Nautical_Twilight object
Astronomical_Twilight object
Year              float64
Month             object
Day               float64
Hour              float64
Weekday           object
```

```
Time_Duration(min)      float64  
dtypes: datetime64[ns](2), float64(32), object(21)  
memory usage: 958.7+ MB
```

```
In [6]: # Double check to make sure no more negative td  
df[td][df[td]<=0]
```

```
Out[6]: Series([], Name: Time_Duration(min), dtype: float64)
```

Step 4. Deal with outliers

B. Fill outliers with median values

```
In [7]: # Remove outliers for Time_Duration(min): n * standard_deviation (n=3), backfill with median

n=3

median = df[td].median()
std = df[td].std()
outliers = (df[td] - median).abs() > std*n

# Set outliers to NAN
df[outliers] = np.nan

# Fill NAN with median
df[td].fillna(median, inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2243926 entries, 0 to 2243938
Data columns (total 55 columns):
ID                object
Source            object
TMC               float64
Severity          float64
Start_Time        datetime64[ns]
End_Time          datetime64[ns]
Start_Lat         float64
Start_Lng         float64
End_Lat           float64
End_Lng           float64
Distance(mi)      float64
Description        object
Number            float64
Street            object
Side              object
City              object
County            object
State             object
Zipcode           object
Country           object
Timezone          object
Airport_Code      object
Weather_Stamp     object
Temperature(F)    float64
Wind_Chill(F)     float64
Humidity(%)       float64
Pressure(in)      float64
Visibility(mi)    float64
Wind_Direction    object
Wind_Speed(mph)   float64
Precipitation(in) float64
Weather_Condition object
Amenity           float64
Bump              float64
Crossing          float64
Give_Way          float64
Junction          float64
No_Exit           float64
Railway           float64
Roundabout        float64
Station           float64
Stop              float64
Traffic_Calming   float64
Traffic_Signal    float64
Turning_Loop      float64
Sunrise_Sunset    object
Civil_Twilight    object
Nautical_Twilight object
Astronomical_Twilight object
Year              float64
Month             object
Day               float64
Hour              float64
Weekday           object
```

```
Time_Duration(min)          float64
dtypes: datetime64[ns](2), float64(32), object(21)
memory usage: 958.7+ MB
```

```
In [8]: # Print time_duration information
print('Max time to clear an accident: {} minutes or {} hours or {} days; Min t
o clear an accident td: {} minutes.'.format(df[td].max(),round(df[td].max()/60
), round(df[td].max()/60/24), df[td].min()))
```

Max time to clear an accident: 12424.0 minutes or 207 hours or 9 days; Min to clear an accident td: 1.0 minutes.

```
In [9]: # Export the data
# df.to_csv('./US_Accidents_May19_clean.csv',index=False)
```

Step 5. Select a list of features for machine learning algorithms

Only select relevant columns without overwhelming the computer

```
In [10]: # Set the list of features to include in Machine Learning
feature_lst=['Source','TMC','Severity','Start_Lng','Start_Lat','Distance(mi)',
'Side','City','County','State','Timezone','Temperature(F)','Humidity(%)','Pres
sure(in)', 'Visibility(mi)', 'Wind_Direction','Weather_Condition','Amenity','B
ump','Crossing','Give_Way','Junction','No_Exit','Railway','Roundabout','Statio
n','Stop','Traffic_Calming','Traffic_Signal','Turning_Loop','Sunrise_Sunset',
'Hour','Weekday', 'Time_Duration(min)']
```

In [11]: *# Select the dataset to include only the selected features*

```
df_sel=df[feature_lst].copy()
df_sel.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2243926 entries, 0 to 2243938
Data columns (total 34 columns):
Source                object
TMC                   float64
Severity              float64
Start_Lng             float64
Start_Lat             float64
Distance(mi)          float64
Side                  object
City                  object
County                object
State                 object
Timezone              object
Temperature(F)         float64
Humidity(%)           float64
Pressure(in)          float64
Visibility(mi)         float64
Wind_Direction         object
Weather_Condition      object
Amenity               float64
Bump                  float64
Crossing              float64
Give_Way              float64
Junction              float64
No_Exit               float64
Railway               float64
Roundabout            float64
Station               float64
Stop                  float64
Traffic_Calming        float64
Traffic_Signal         float64
Turning_Loop          float64
Sunrise_Sunset         object
Hour                  float64
Weekday                object
Time_Duration(min)     float64
dtypes: float64(24), object(10)
memory usage: 599.2+ MB
```

In [12]: *# Export the data with selected features*

```
# df_sel.to_csv('./US_Accidents_May19_clean_sel.csv', index=False)
```

Step 6. Drop rows with missing values

```
In [13]: # Check missing values
df_sel.isnull().mean()
```

```
Out[13]: Source                0.000217
TMC                0.230302
Severity           0.000217
Start_Lng          0.000217
Start_Lat          0.000217
Distance(mi)       0.000217
Side               0.000217
City               0.000248
County             0.000217
State              0.000217
Timezone           0.001172
Temperature(F)     0.027943
Humidity(%)        0.028925
Pressure(in)       0.025724
Visibility(mi)     0.031997
Wind_Direction     0.021230
Weather_Condition  0.032285
Amenity            0.000217
Bump               0.000217
Crossing           0.000217
Give_Way           0.000217
Junction           0.000217
No_Exit            0.000217
Railway            0.000217
Roundabout        0.000217
Station            0.000217
Stop               0.000217
Traffic_Calming    0.000217
Traffic_Signal     0.000217
Turning_Loop       0.000217
Sunrise_Sunset     0.000252
Hour               0.000217
Weekday            0.000217
Time_Duration(min) 0.000000
dtype: float64
```

```
In [14]: df_sel.dropna(subset=df_sel.columns[df_sel.isnull().mean()!=0], how='any', axis=0, inplace=True)
df_sel.shape
```

```
Out[14]: (1663631, 34)
```

```
In [15]: # Export the data with selected features
# df_sel.to_csv('./US_Accidents_May19_clean_sel_dropna.csv', index=False)
```

Step 7. Select the state of interest: CA

Due to the limitation of personal laptop, the whole US dataset is too big to handle


```
In [16]: # Import data if it was already exported based on previous work  
# df_sel=pd.read_csv('./US_Accidents_May19_clean_sel_dropna.csv')
```

```
In [17]: # Set state
state='CA'

# Select the state of Pennsylvania
df_state=df_sel.loc[df_sel.State==state]
df_state.drop('State',axis=1, inplace=True)
df_state.info()
```

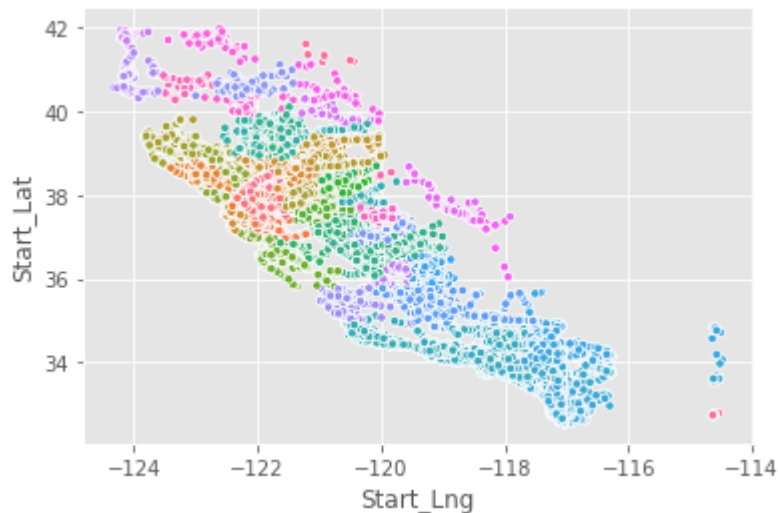
```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 354854 entries, 728 to 1727176
Data columns (total 33 columns):
Source                354854 non-null object
TMC                   354854 non-null float64
Severity              354854 non-null float64
Start_Lng             354854 non-null float64
Start_Lat             354854 non-null float64
Distance(mi)          354854 non-null float64
Side                  354854 non-null object
City                  354854 non-null object
County                354854 non-null object
Timezone              354854 non-null object
Temperature(F)         354854 non-null float64
Humidity(%)            354854 non-null float64
Pressure(in)           354854 non-null float64
Visibility(mi)         354854 non-null float64
Wind_Direction         354854 non-null object
Weather_Condition      354854 non-null object
Amenity               354854 non-null float64
Bump                   354854 non-null float64
Crossing              354854 non-null float64
Give_Way              354854 non-null float64
Junction              354854 non-null float64
No_Exit               354854 non-null float64
Railway               354854 non-null float64
Roundabout            354854 non-null float64
Station               354854 non-null float64
Stop                  354854 non-null float64
Traffic_Calming        354854 non-null float64
Traffic_Signal         354854 non-null float64
Turning_Loop          354854 non-null float64
Sunrise_Sunset        354854 non-null object
Hour                  354854 non-null float64
Weekday               354854 non-null object
Time_Duration(min)     354854 non-null float64
dtypes: float64(24), object(9)
memory usage: 92.0+ MB
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\frame.py:4102: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
errors=errors,

```
In [18]: # Map of accidents, color code by county

sns.scatterplot(x='Start_Lng', y='Start_Lat', data=df_state, hue='County', legend=False, s=20)
plt.show()
```



Step 8. Deal with categorical data: pd.get_dummies()

```
In [19]: # Generate dummies for categorical data
df_state_dummy = pd.get_dummies(df_state, drop_first=True)

# Export data
df_state_dummy.to_csv('./US_Accidents_May19_{state}_dummy.csv'.format(state), index=False)

df_state_dummy.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 354854 entries, 728 to 1727176
Columns: 1162 entries, TMC to Weekday_Wed
dtypes: float64(24), uint8(1138)
memory usage: 452.8 MB
```

Step 9. Predict the accident severity with various supervised machine learning algorithms

Data preparation: train_test_split

```
In [20]: # Assign the data
df=df_state_dummy

# Set the target for the prediction
target='Severity'

# Create arrays for the features and the response variable

# set X and y
y = df[target]
X = df.drop(target, axis=1)

# Split the data set into training and testing data sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=21, stratify=y)
```

```
In [21]: # List of classification algorithms
algo_lst=['Logistic Regression', 'K-Nearest Neighbors', 'Decision Trees', 'Random Forest']

# Initialize an empty list for the accuracy for each algorithm
accuracy_lst=[]
```

Step 9. Predict the accident severity with various supervised machine learning algorithms

Algorithm A. Logistic regression

```
In [22]: # Logistic regression
lr = LogisticRegression(random_state=0)
lr.fit(X_train,y_train)
y_pred=lr.predict(X_test)

# Get the accuracy score
acc=accuracy_score(y_test, y_pred)

# Append to the accuracy list
accuracy_lst.append(acc)

print("[Logistic regression algorithm] accuracy_score: {:.3f}.".format(acc))
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
FutureWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:469: FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify the multi_class option to silence this warning.
"this warning.", FutureWarning)

[Logistic regression algorithm] accuracy_score: 0.809.

Step 9. Predict the accident severity with various supervised machine learning algorithms

Algorithm B. The K-Nearest Neighbors (KNN) algorithm

KNN with 6 neighbors

```
In [23]: # Create a k-NN classifier with 6 neighbors
knn = KNeighborsClassifier(n_neighbors=6)

# Fit the classifier to the data
knn.fit(X_train,y_train)

# Predict the labels for the training data X
y_pred = knn.predict(X_test)

# Get the accuracy score
acc=accuracy_score(y_test, y_pred)

# Append to the accuracy list
accuracy_lst.append(acc)

print('[K-Nearest Neighbors (KNN)] knn.score: {:.3f}.'.format(knn.score(X_test, y_test)))
print('[K-Nearest Neighbors (KNN)] accuracy_score: {:.3f}.'.format(acc))
```

[K-Nearest Neighbors (KNN)] knn.score: 0.604.
[K-Nearest Neighbors (KNN)] accuracy_score: 0.604.

Step 9. Predict the accident severity with various supervised machine learning algorithms

Algorithm B. The K-Nearest Neighbors (KNN) algorithm

Optimize the number of neighbors: plot the accuracy versus number of neighbors

It took too much time, skip this part.

Setup arrays to store train and test accuracies

```
neighbors = np.arange(1, 9) train_accuracy = np.empty(len(neighbors)) test_accuracy =  
np.empty(len(neighbors))
```

Loop over different values of k

```
for i, n_neighbor in enumerate(neighbors):  
  
    # Setup a k-NN Classifier with n_neighbor  
    knn = KNeighborsClassifier(n_neighbors=n_neighbor)  
  
    # Fit the classifier to the training data  
    knn.fit(X_train,y_train)  
  
    #Compute accuracy on the training set  
    train_accuracy[i] = knn.score(X_train, y_train)  
  
    #Compute accuracy on the testing set  
    test_accuracy[i] = knn.score(X_test, y_test)
```

Generate plot

```
plt.title('k-NN: Varying Number of Neighbors') plt.plot(neighbors, test_accuracy, label = 'Testing Accuracy')  
plt.plot(neighbors, train_accuracy, label = 'Training Accuracy') plt.legend() plt.xlabel('Number of Neighbors')  
plt.ylabel('Accuracy') plt.show()
```

Step 9. Predict the accident severity with various supervised machine learning algorithms

Algorithm C. Decision Tree

```
In [24]: # Decision tree algorithm

# Instantiate dt_entropy, set 'entropy' as the information criterion
dt_entropy = DecisionTreeClassifier(max_depth=8, criterion='entropy', random_s
tate=1)

# Fit dt_entropy to the training set
dt_entropy.fit(X_train, y_train)

# Use dt_entropy to predict test set labels
y_pred= dt_entropy.predict(X_test)

# Evaluate accuracy_entropy
accuracy_entropy = accuracy_score(y_test, y_pred)

# Print accuracy_entropy
print('[Decision Tree -- entropy] accuracy_score: {:.3f}'.format(accuracy_ent
ropy))

# Instantiate dt_gini, set 'gini' as the information criterion
dt_gini = DecisionTreeClassifier(max_depth=8, criterion='gini', random_state=1
)

# Fit dt_entropy to the training set
dt_gini.fit(X_train, y_train)

# Use dt_entropy to predict test set labels
y_pred= dt_gini.predict(X_test)

# Evaluate accuracy_entropy
accuracy_gini = accuracy_score(y_test, y_pred)

# Append to the accuracy List
acc=accuracy_gini
accuracy_lst.append(acc)

# Print accuracy_gini
print('[Decision Tree -- gini] accuracy_score: {:.3f}'.format(accuracy_gini))

[Decision Tree -- entropy] accuracy_score: 0.743.
[Decision Tree -- gini] accuracy_score: 0.751.
```

Step 9. Predict the accident severity with various supervised machine learning algorithms

Algorithm D. Random Forest

n_estimators=100

```
In [25]: # Random Forest algorithm

#Create a Gaussian Classifier
clf=RandomForestClassifier(n_estimators=100)

#Train the model using the training sets y_pred=clf.predict(X_test)
clf.fit(X_train,y_train)

y_pred=clf.predict(X_test)

# Get the accuracy score
acc=accuracy_score(y_test, y_pred)

# Append to the accuracy list
accuracy_lst.append(acc)

# Model Accuracy, how often is the classifier correct?
print("[Random forest algorithm] accuracy_score: {:.3f}.".format(acc))

[Random forest algorithm] accuracy_score: 0.892.
```

Step 9. Predict the accident severity with various supervised machine learning algorithms


Algorithm D. Random Forest

Visualize important features


```
In [26]: feature_imp = pd.Series(clf.feature_importances_, index=X.columns).sort_values(
        ascending=False)

# Creating a bar plot, displaying only the top k features
k=10
sns.barplot(x=feature_imp[:10], y=feature_imp.index[:k])
# Add labels to your graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()
plt.show()
```

No handles with labels found to put in legend.



```
In [27]: # List top k important features
k=20
feature_imp.sort_values(ascending=False)[:k]
```

```
Out[27]: Start_Lat      0.134451
Start_Lng      0.134298
Humidity(%)    0.042199
Pressure(in)   0.042120
Temperature(F) 0.041817
Hour           0.037062
Side_L         0.032037
Side_R         0.029736
Time_Duration(min) 0.026945
Traffic_Signal 0.021809
Distance(mi)   0.014331
Visibility(mi) 0.012959
TMC            0.010856
Junction       0.007621
Sunrise_Sunset_Night 0.007340
County_Los Angeles 0.007255
County_Sonoma  0.005987
City_Oakland   0.005886
Weather_Condition_Clear 0.005804
Weekday_Wed    0.005721
dtype: float64
```

Step 9. Predict the accident severity with various supervised machine learning algorithms

Algorithm D. Random Forest

Select the top important features, set the threshold

```
In [28]: # Create a selector object that will use the random forest classifier to identify
# features that have an importance of more than 0.03
sfm = SelectFromModel(clf, threshold=0.03)

# Train the selector
sfm.fit(X_train, y_train)

feat_labels=X.columns

# Print the names of the most important features
for feature_list_index in sfm.get_support(indices=True):
    print(feat_labels[feature_list_index])

Start_Lng
Start_Lat
Temperature(F)
Humidity(%)
Pressure(in)
Hour
Side_R
```

```
In [29]: # Transform the data to create a new dataset containing only the most important features
# Note: We have to apply the transform to both the training X and test X data.
X_important_train = sfm.transform(X_train)
X_important_test = sfm.transform(X_test)

# Create a new random forest classifier for the most important features
clf_important = RandomForestClassifier(n_estimators=100, random_state=0, n_jobs=-1)

# Train the new classifier on the new dataset containing the most important features
clf_important.fit(X_important_train, y_train)
```

```
Out[29]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=-1, oob_score=False, random_state=0, verbose=0,
                                warm_start=False)
```

```
In [30]: # Apply The Full Featured Classifier To The Test Data
y_pred = clf.predict(X_test)

# View The Accuracy Of Our Full Feature Model
print('[Random forest algorithm -- Full feature] accuracy_score: {:.3f}.'.format(accuracy_score(y_test, y_pred)))

# Apply The Full Featured Classifier To The Test Data
y_important_pred = clf_important.predict(X_important_test)

# View The Accuracy Of Our Limited Feature Model
print('[Random forest algorithm -- Limited feature] accuracy_score: {:.3f}.'.format(accuracy_score(y_test, y_important_pred)))

[Random forest algorithm -- Full feature] accuracy_score: 0.892.
[Random forest algorithm -- Limited feature] accuracy_score: 0.911.
```

Step 9. Predict the accident severity with various supervised machine learning algorithms

Plot the accuracy score versus algorithm

```

In [31]: # Make a plot of the accuracy scores for different algorithms

# Generate a list of ticks for y-axis
y_ticks=np.arange(len(algo_lst))

# Combine the list of algorithms and list of accuracy scores into a dataframe,
sort the value based on accuracy score
df_acc=pd.DataFrame(list(zip(algo_lst, accuracy_lst)), columns=['Algorithm','A
ccuracy_Score']).sort_values(by=['Accuracy_Score'],ascending = True)

# Export to a file
df_acc.to_csv('./Accuracy_scores_algorithms_{}.csv'.format(state),index=False)

# Make a plot
ax=df_acc.plot.barh('Algorithm', 'Accuracy_Score', align='center',legend=False
,color='0.5')

# Add the data label on to the plot
for i in ax.patches:
    # get_width pulls left or right; get_y pushes up or down
    ax.text(i.get_width()+0.02, i.get_y()+0.2, str(round(i.get_width(),2)), fo
ntsize=10)

# Set the limit, lables, ticks and title
plt.xlim(0,1.05)
plt.xlabel('Accuracy Score')
plt.yticks(y_ticks, df_acc['Algorithm'], rotation=0)
plt.title('[] Which algorithm is better?'.format(state))

plt.show()

```

